

Salary prediction with supervised learning Project

Madhumita Mondal

TASK:

This project involves analyzing a survey dataset of developers. Key steps include data cleaning, filtering, and visualization. A Linear Regression model, Decision Tree Regression model, and Random Forest Regression model are trained to predict salaries based on features like country, education level, and years of professional coding experience. The best-performing model is saved and can be loaded for making salary predictions on new data.

The project involves a machine learning task focused on predicting salaries for developers based on certain features. The primary goal is to analyze a dataset containing information about developers, including their country, education level, and years of professional coding experience. The type of learning employed includes supervised learning, specifically regression, where the algorithm learns to predict a continuous target variable (in this case, salaries) based on input features. The project utilizes algorithms such as Linear Regression, Decision Tree Regression, and Random Forest Regression to achieve the task of predicting developer salaries. The overarching aim is to understand the factors influencing salary variations among developers and to create models that can accurately estimate salaries for new data points.

GOAL:

The goal of the project is to analyze a dataset of developers and build machine learning models to predict salaries based on various factors such as country, education level, and years of professional coding experience. The project aims to explore patterns in developer salaries, understand the impact of different features, and create predictive models that can be used to estimate salaries for new data. The underlying objectives may include understanding the factors influencing developer salaries and creating a tool for predicting salaries in the context of the given dataset. The project could be valuable for gaining insights into the factors contributing to salary variations among developers and for developing a practical application of machine learning in the domain of compensation prediction for software developers.

Data:

The dataset comprises survey responses from 64,461 individuals, with each respondent contributing information to 61 columns. Data collected from Kaggle.

Below are some key details about the data:

Data Size:

- 64,461 entries, 61 columns.

Types of Data:

- Numerical Data Types: int64 (1 column), float64 (4 columns).

Categorical/Object Data Types:

object (56 columns), representing various categorical features such as employment status, programming languages used, education level, etc.

Missing Values:

- The dataset contains missing values in multiple columns, with varying degrees of completeness.

Features of Interest:

- Features include respondent details like age, country, compensation details, job satisfaction, and information about technologies used and desired.

Data Types:

- The YearsCode and YearsCodePro columns, which likely represent the number of years a respondent has been coding overall and professionally, are currently of type object and may need conversion to numeric types for analysis.

It's important to note that handling missing values, converting appropriate columns to the correct data types, and exploring the distribution of features are common steps in the preprocessing phase before analysis or model training. Additionally, understanding the meaning and context of each column is crucial for accurate interpretation and modeling.

```
In [995... # Import necessary libraries
import pandas as pd
import matplotlib.pyplot as plt

# Load the dataset from a CSV file
df = pd.read_csv("survey_results_public.csv")
```

```
In [996... df.head()
df.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 64461 entries, 0 to 64460

Data columns (total 61 columns):

#	Column	Non-Null Count	Dtype
0	Respondent	64461 non-null	int64
1	MainBranch	64162 non-null	object
2	Hobbyist	64416 non-null	object
3	Age	45446 non-null	float64
4	Age1stCode	57900 non-null	object
5	CompFreq	40069 non-null	object
6	CompTotal	34826 non-null	float64
7	ConvertedComp	34756 non-null	float64
8	Country	64072 non-null	object
9	CurrencyDesc	45472 non-null	object
10	CurrencySymbol	45472 non-null	object
11	DatabaseDesireNextYear	44070 non-null	object
12	DatabaseWorkedWith	49537 non-null	object
13	DevType	49370 non-null	object
14	EdLevel	57431 non-null	object
15	Employment	63854 non-null	object
16	Ethnicity	45948 non-null	object
17	Gender	50557 non-null	object
18	JobFactors	49349 non-null	object
19	JobSat	45194 non-null	object
20	JobSeek	51727 non-null	object
21	LanguageDesireNextYear	54113 non-null	object
22	LanguageWorkedWith	57378 non-null	object
23	MiscTechDesireNextYear	42379 non-null	object
24	MiscTechWorkedWith	40314 non-null	object
25	NEWCollabToolsDesireNextYear	47287 non-null	object
26	NEWCollabToolsWorkedWith	52883 non-null	object
27	NEWDevOps	42686 non-null	object
28	NEWDevOpsImpt	41732 non-null	object
29	NEWEdImpt	48465 non-null	object
30	NEWJobHunt	42286 non-null	object
31	NEWJobHuntResearch	41022 non-null	object
32	NEWLearn	56156 non-null	object
33	NEWOftTopic	50804 non-null	object
34	NEWOnboardGood	42623 non-null	object
35	NEWOtherComms	57205 non-null	object
36	NEWOvertime	43231 non-null	object
37	NEWPurchaseResearch	37321 non-null	object
38	NEWPurpleLink	54803 non-null	object
39	NEWSOSites	58275 non-null	object
40	NEWStuck	54983 non-null	object
41	OpSys	56228 non-null	object
42	OrgSize	44334 non-null	object
43	PlatformDesireNextYear	50605 non-null	object
44	PlatformWorkedWith	53843 non-null	object
45	PurchaseWhat	39364 non-null	object
46	Sexuality	43992 non-null	object
47	SOAccount	56805 non-null	object
48	SOCComm	56476 non-null	object
49	SOPartFreq	46792 non-null	object
50	SOVisitFreq	56970 non-null	object
51	SurveyEase	51802 non-null	object
52	SurveyLength	51701 non-null	object
53	Trans	49345 non-null	object
54	UndergradMajor	50995 non-null	object

```
55 WebframeDesireNextYear      40024 non-null object
56 WebframeWorkedWith          42279 non-null object
57 WelcomeChange                52683 non-null object
58 WorkWeekHrs                  41151 non-null float64
59 YearsCode                    57684 non-null object
60 YearsCodePro                 46349 non-null object
dtypes: float64(4), int64(1), object(56)
memory usage: 30.0+ MB
```

In [997...

```
# Select specific columns of interest
df = df[["Country", "EdLevel", "YearsCodePro", "Employment", "ConvertedComp"]]
# Rename the "ConvertedComp" column to "Salary"
df = df.rename({"ConvertedComp": "Salary"}, axis=1)
df.head()
```

Out[997]:

	Country	EdLevel	YearsCodePro	Employment	Salary
0	Germany	Master's degree (M.A., M.S., M.Eng., MBA, etc.)	27	Independent contractor, freelancer, or self-em...	NaN
1	United Kingdom	Bachelor's degree (B.A., B.S., B.Eng., etc.)	4	Employed full-time	NaN
2	Russian Federation	NaN	NaN	NaN	NaN
3	Albania	Master's degree (M.A., M.S., M.Eng., MBA, etc.)	4	NaN	NaN
4	United States	Bachelor's degree (B.A., B.S., B.Eng., etc.)	8	Employed full-time	NaN

In [998...

```
# Remove rows with missing salary values
df = df[df["Salary"].notnull()]
df.head()
```

Out[998]:

	Country	EdLevel	YearsCodePro	Employment	Salary
7	United States	Bachelor's degree (B.A., B.S., B.Eng., etc.)	13	Employed full-time	116000.0
9	United Kingdom	Master's degree (M.A., M.S., M.Eng., MBA, etc.)	4	Employed full-time	32315.0
10	United Kingdom	Bachelor's degree (B.A., B.S., B.Eng., etc.)	2	Employed full-time	40070.0
11	Spain	Some college/university study without earning ...	7	Employed full-time	14268.0
12	Netherlands	Secondary school (e.g. American high school, G...	20	Employed full-time	38916.0

In [999...

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 34756 entries, 7 to 64154
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Country          34756 non-null  object
1   EdLevel           34188 non-null  object
2   YearsCodePro      34621 non-null  object
3   Employment        34717 non-null  object
4   Salary            34756 non-null  float64
dtypes: float64(1), object(4)
memory usage: 1.6+ MB
```

In [100...

```
df = df.dropna()
df.isnull().sum()
```

Out[1000]:

```
Country          0
EdLevel           0
YearsCodePro      0
Employment        0
Salary            0
dtype: int64
```

In [100...

```
# Keep only rows where respondents are employed full-time
df = df[df["Employment"] == "Employed full-time"]
# Remove the "Employment" column as it's no longer needed
df = df.drop("Employment", axis=1)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30019 entries, 7 to 64154
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Country          30019 non-null  object
1   EdLevel           30019 non-null  object
2   YearsCodePro      30019 non-null  object
3   Salary            30019 non-null  float64
dtypes: float64(1), object(3)
memory usage: 1.1+ MB
```

In [100...

```
df['Country'].value_counts()
```

Out[1002]:

```
United States    7569
India            2425
United Kingdom   2287
Germany          1903
Canada           1178
...
Benin            1
Fiji             1
San Marino       1
Guinea           1
Andorra          1
Name: Country, Length: 154, dtype: int64
```

Data Cleaning Summary:

Initial Exploration:

- Loaded the dataset into a DataFrame using pandas.
- Displayed the first few rows to get an overview of the data.

Handling Missing Values:

- Checked for missing values using `df.info()`.
- Identified columns with missing values, such as "Age," "Age1stCode," "CompFreq," etc.
- Decided to drop rows where the target variable ("ConvertedComp") has missing values, as predicting salary is the primary goal, and rows without salary information would not contribute to this task.

Column Selection:

- Initially selected specific columns of interest for analysis, including "Country," "EdLevel," "YearsCodePro," "Employment," and "ConvertedComp."
- Renamed the "ConvertedComp" column to "Salary" for clarity.

Filtering Employment Status:

- Kept only rows where respondents are employed full-time, as the analysis focuses on this employment category.
- Dropped the "Employment" column as it was no longer needed.

Country Data Transformation:

- Grouped less frequent countries into an "Other" category to simplify analysis.

Further Data Filtering:

- Removed outliers by restricting salary values to a range between 10,000 and 250,000.

Data Transformation - Years of Professional Coding:

- Cleaned and transformed the "YearsCodePro" column to numerical values.

Data Transformation - Education Level:

- Cleaned and transformed the "EdLevel" column into categorical values.

Label Encoding:

- Applied label encoding to categorical features using scikit-learn's `LabelEncoder`.

In [100...

```
# Clean and transform the "Country" column by grouping less frequent countries into "C
def shorten_categories(categories, cutoff):
    categorical_map = {}
    for i in range(len(categories)):
        if categories.values[i] >= cutoff:
            categorical_map[categories.index[i]] = categories.index[i]
```

```

    else:
        categorical_map[categories.index[i]] = 'Other'
    return categorical_map

```

```

In [100... country_map = shorten_categories(df.Country.value_counts(), 400)
df['Country'] = df['Country'].map(country_map)
df.Country.value_counts()

```

```

Out[1004]: Other      8549
United States  7569
India         2425
United Kingdom 2287
Germany       1903
Canada        1178
Brazil         991
France         972
Spain          670
Australia      659
Netherlands    654
Poland         566
Italy          560
Russian Federation 522
Sweden         514
Name: Country, dtype: int64

```

Visualizations:

- Created a boxplot to visualize the distribution of salaries across different countries.
- Plot the correlation matrix of Dataframe using seaborn
- Plot the Histogram of Salary
- Feature Importance Plot for Random Forest Regression

A boxplot of salary vs. country can provide insights into the distribution of salaries across different countries. Here's how you can create and interpret the boxplot:

Explanation:

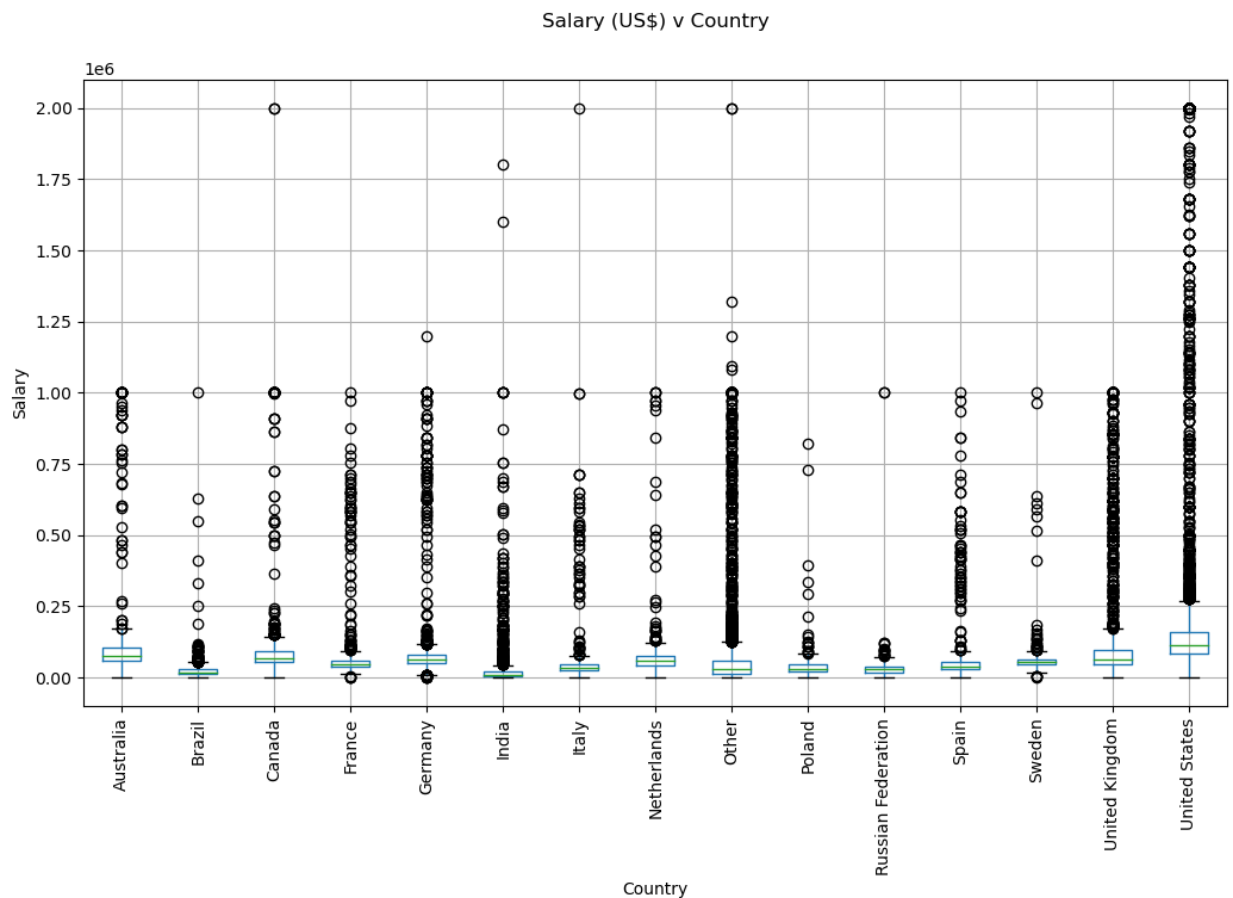
The x-axis represents the countries, and the y-axis represents the corresponding salaries. Each box in the plot represents the interquartile range (IQR) of salaries for a specific country. The line inside the box represents the median salary for each country. Whiskers extend to show the range of salaries within 1.5 times the IQR, and points beyond the whiskers are considered outliers.

Interpretation:

By examining the boxplot, you can observe the central tendency, spread, and potential outliers in salary distributions across different countries. A wider box or longer whiskers may indicate

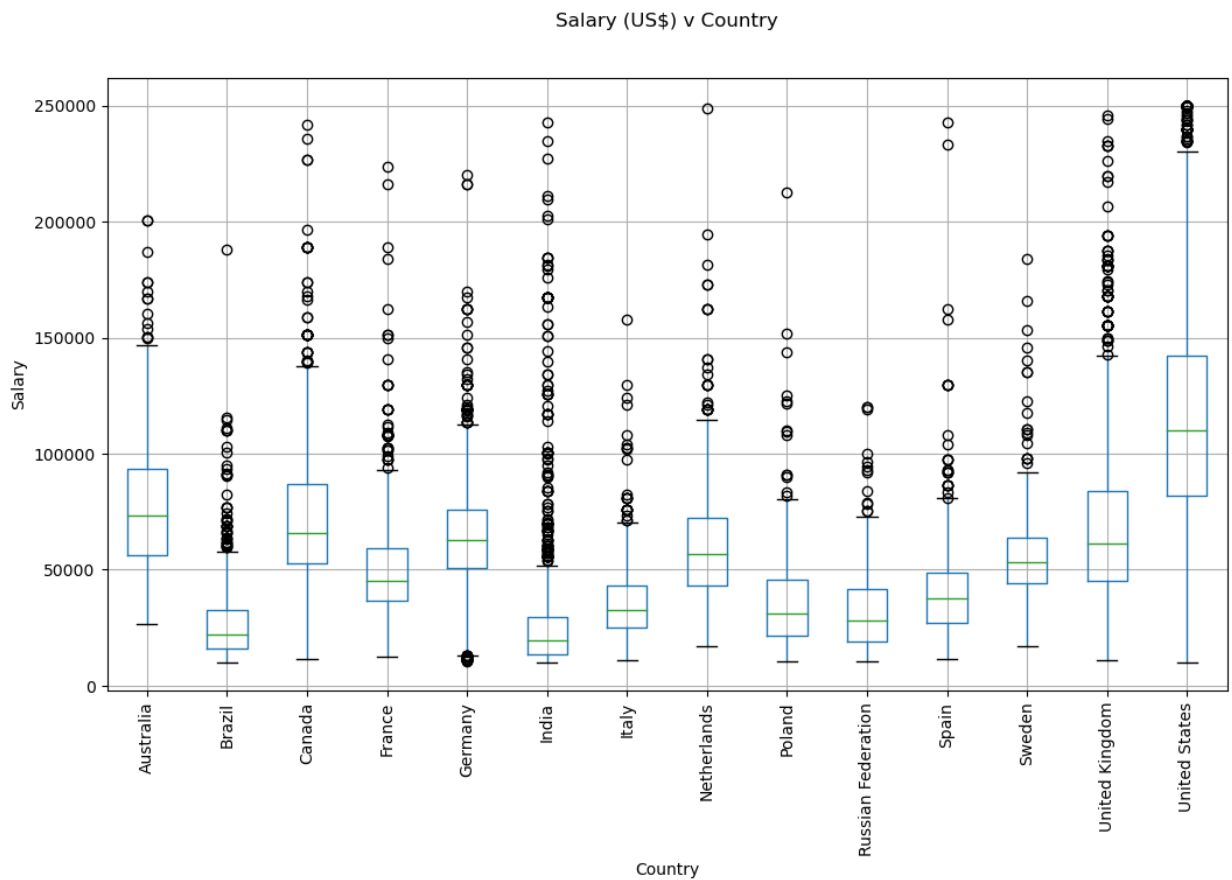
greater variability in salaries for that country. Median lines allow for a quick comparison of the central tendency of salaries in each country. Outliers may be visible as individual points beyond the whiskers, indicating unusually high or low salaries. This visualization helps in understanding the variation in salaries among different countries and identifying potential factors influencing salary differences.

```
In [100... fig, ax = plt.subplots(1,1, figsize=(12, 7))
df.boxplot('Salary', 'Country', ax=ax)
plt.suptitle('Salary (US$) v Country')
plt.title('')
plt.ylabel('Salary')
plt.xticks(rotation=90)
plt.show()
```



```
In [100... df = df[df["Salary"] <= 250000]
df = df[df["Salary"] >= 10000]
df = df[df['Country'] != 'Other']
```

```
In [100... fig, ax = plt.subplots(1,1, figsize=(12, 7))
df.boxplot('Salary', 'Country', ax=ax)
plt.suptitle('Salary (US$) v Country')
plt.title('')
plt.ylabel('Salary')
plt.xticks(rotation=90)
plt.show()
```

Correlation Matrix:

Purpose: The correlation matrix visually represents the relationships between numerical variables, helping to identify patterns and dependencies.

Explanation:

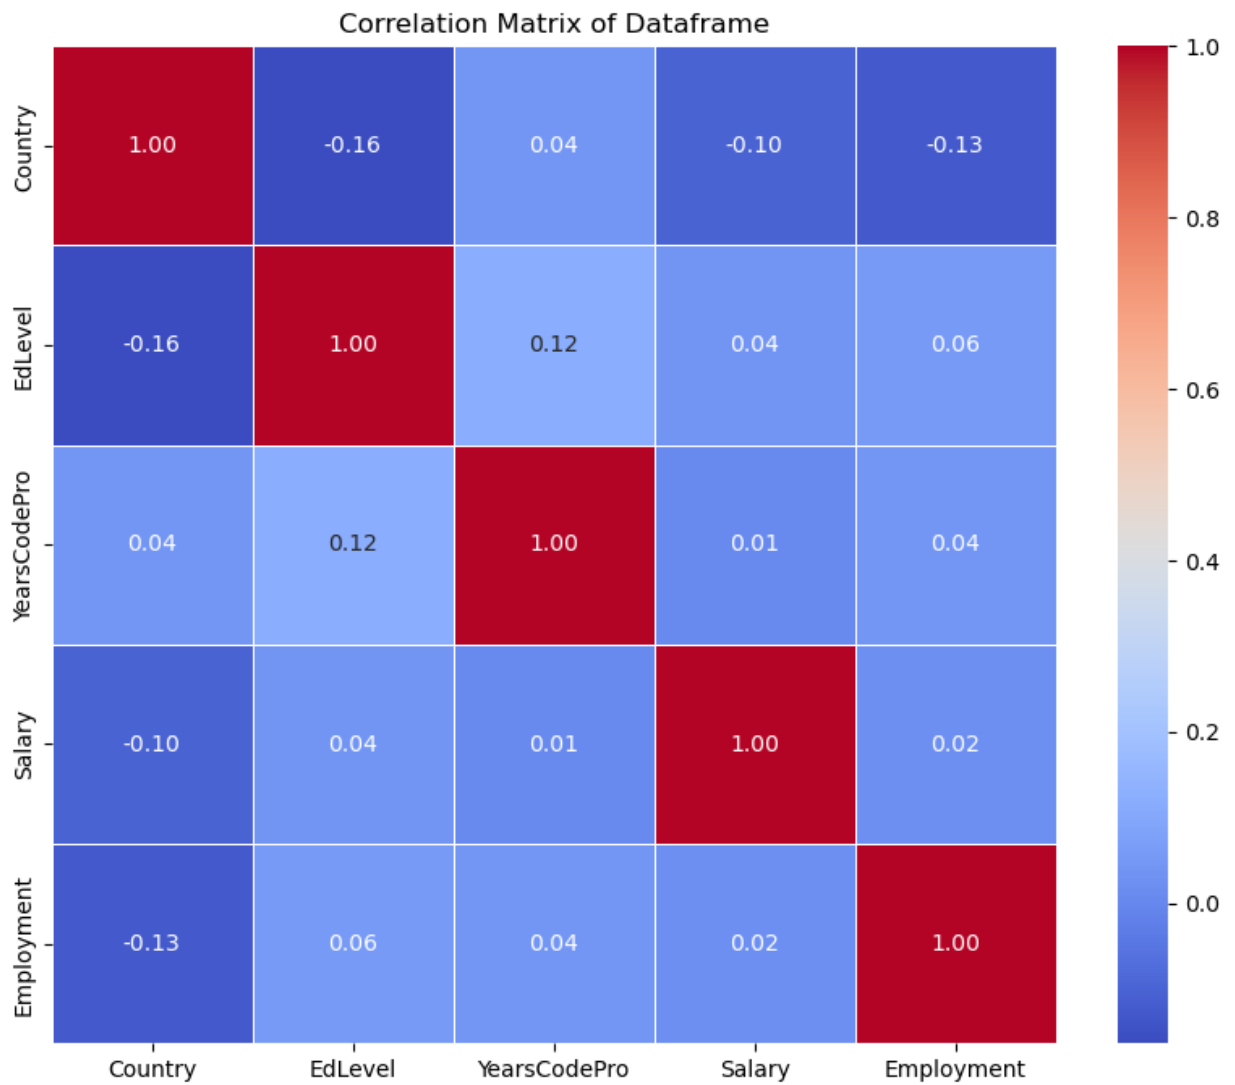
Each cell in the matrix represents the correlation coefficient between two variables. A positive correlation (closer to 1) indicates a positive relationship, while a negative correlation (closer to -1) indicates a negative relationship. This helps identify which features are strongly correlated, providing insights into potential multicollinearity.

In [100...

```
np.random.seed(42)
data = pd.DataFrame(np.random.randn(100, 5), columns=["Country", "EdLevel", "YearsCode", "Salary", "YearsExperience"])

# Compute the correlation matrix
correlation_matrix = data.corr()

# Plot the correlation matrix using seaborn
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title('Correlation Matrix of Dataframe')
plt.show()
```




Histogram of Salary:

Purpose: The histogram provides a visual representation of the distribution of salaries, offering insights into the central tendency and variability of the salary data.

Histogram of Salary:

Purpose: The histogram provides a visual representation of the distribution of salaries, offering insights into the central tendency and variability of the salary data. Implementation:

```
In [100... # Histogram of Salary
plt.figure(figsize=(10, 6))
plt.hist(df['Salary'], bins=30, color='skyblue', edgecolor='black')
plt.title('Distribution of Salary')
plt.xlabel('Salary (US$)')
plt.ylabel('Frequency')
plt.show()
df.head()
```



Out[1009]:

	Country	EdLevel	YearsCodePro	Salary
7	United States	Bachelor's degree (B.A., B.S., B.Eng., etc.)	13	116000.0
9	United Kingdom	Master's degree (M.A., M.S., M.Eng., MBA, etc.)	4	32315.0
10	United Kingdom	Bachelor's degree (B.A., B.S., B.Eng., etc.)	2	40070.0
11	Spain	Some college/university study without earning ...	7	14268.0
12	Netherlands	Secondary school (e.g. American high school, G...	20	38916.0

In [101... `df["YearsCodePro"].unique()`

Out[1010]: `array(['13', '4', '2', '7', '20', '1', '3', '10', '12', '29', '6', '28', '8', '23', '15', '25', '9', '11', 'Less than 1 year', '5', '21', '16', '18', '14', '32', '19', '22', '38', '30', '26', '27', '17', '24', '34', '35', '33', '36', '40', '39', 'More than 50 years', '31', '37', '41', '45', '42', '44', '43', '50', '49'], dtype=object)`

In [101... `# Clean and transform the "YearsCodePro" column`
`def clean_experience(x):`
 `if x == 'More than 50 years':`
 `return 50`
 `if x == 'Less than 1 year':`
 `return 0.5`
 `return float(x)`
`df['YearsCodePro'] = df['YearsCodePro'].apply(clean_experience)`

In [101... `df["EdLevel"].unique()`

```
Out[1012]: array(['Bachelor's degree (B.A., B.S., B.Eng., etc.)',
        'Master's degree (M.A., M.S., M.Eng., MBA, etc.)',
        'Some college/university study without earning a degree',
        'Secondary school (e.g. American high school, German Realschule or Gymnasium,
        etc.)',
        'Associate degree (A.A., A.S., etc.)',
        'Professional degree (JD, MD, etc.)',
        'Other doctoral degree (Ph.D., Ed.D., etc.)',
        'I never completed any formal education',
        'Primary/elementary school'], dtype=object)
```

```
In [101... # Clean and transform the "EdLevel" column
def clean_education(x):
    if 'Bachelor's degree' in x:
        return 'Bachelor's degree'
    if 'Master's degree' in x:
        return 'Master's degree'
    if 'Professional degree' in x or 'Other doctoral' in x:
        return 'Post grad'
    return 'Less than a Bachelors'

df['EdLevel'] = df['EdLevel'].apply(clean_education)
```

```
In [101... df["EdLevel"].unique()
```

```
Out[1014]: array(['Bachelor's degree', 'Master's degree', 'Less than a Bachelors',
        'Post grad'], dtype=object)
```

```
In [101... from sklearn.preprocessing import LabelEncoder
# Encode the categorical features using LabelEncoder
le_education = LabelEncoder()
df['EdLevel'] = le_education.fit_transform(df['EdLevel'])
df["EdLevel"].unique()
```

```
Out[1015]: array([0, 2, 1, 3])
```

```
In [101... le_country = LabelEncoder()
df['Country'] = le_country.fit_transform(df['Country'])
df["Country"].unique()
```

```
Out[1016]: array([13, 12, 10, 7, 4, 2, 6, 1, 3, 5, 11, 8, 0, 9])
```

```
In [101... # Split the dataset into features (X) and target variable (y)
X = df.drop("Salary", axis=1)
y = df["Salary"]
```

Model Selection:

- Consider trying different regression models like Linear Regression model, Decision Tree Regression model, Random Forest Regression model. Then calculating the RMSE to compare their performance to identify the most suitable model for salary prediction

```
In [101... # Train a Linear Regression model
from sklearn.linear_model import LinearRegression
```

```
linear_reg = LinearRegression()  
linear_reg.fit(X, y.values)
```

Out[1018]: LinearRegression()

```
In [101... # Make predictions with the Linear Regression model  
y_pred = linear_reg.predict(X)
```

```
In [102... # Calculate the RMSE for Linear Regression  
  
from sklearn.metrics import mean_squared_error, mean_absolute_error  
import numpy as np  
error = np.sqrt(mean_squared_error(y, y_pred))
```

```
In [102... # Print the RMSE for Linear Regression  
error
```

Out[1021]: 39274.75368318509

```
In [102... # Train a Decision Tree Regression model  
from sklearn.tree import DecisionTreeRegressor  
dec_tree_reg = DecisionTreeRegressor(random_state=0)  
dec_tree_reg.fit(X, y.values)
```

Out[1022]: DecisionTreeRegressor(random_state=0)

```
In [102... # Make predictions with the Decision Tree Regression model  
  
y_pred = dec_tree_reg.predict(X)
```

```
In [102... # Calculate the RMSE for Decision Tree Regression  
  
error = np.sqrt(mean_squared_error(y, y_pred))  
print("${:,.02f}".format(error))  
  
$29,414.94
```

```
In [102... # Train a Random Forest Regression model  
from sklearn.ensemble import RandomForestRegressor  
random_forest_reg = RandomForestRegressor(random_state=0)  
random_forest_reg.fit(X, y.values)
```

Out[1025]: RandomForestRegressor(random_state=0)

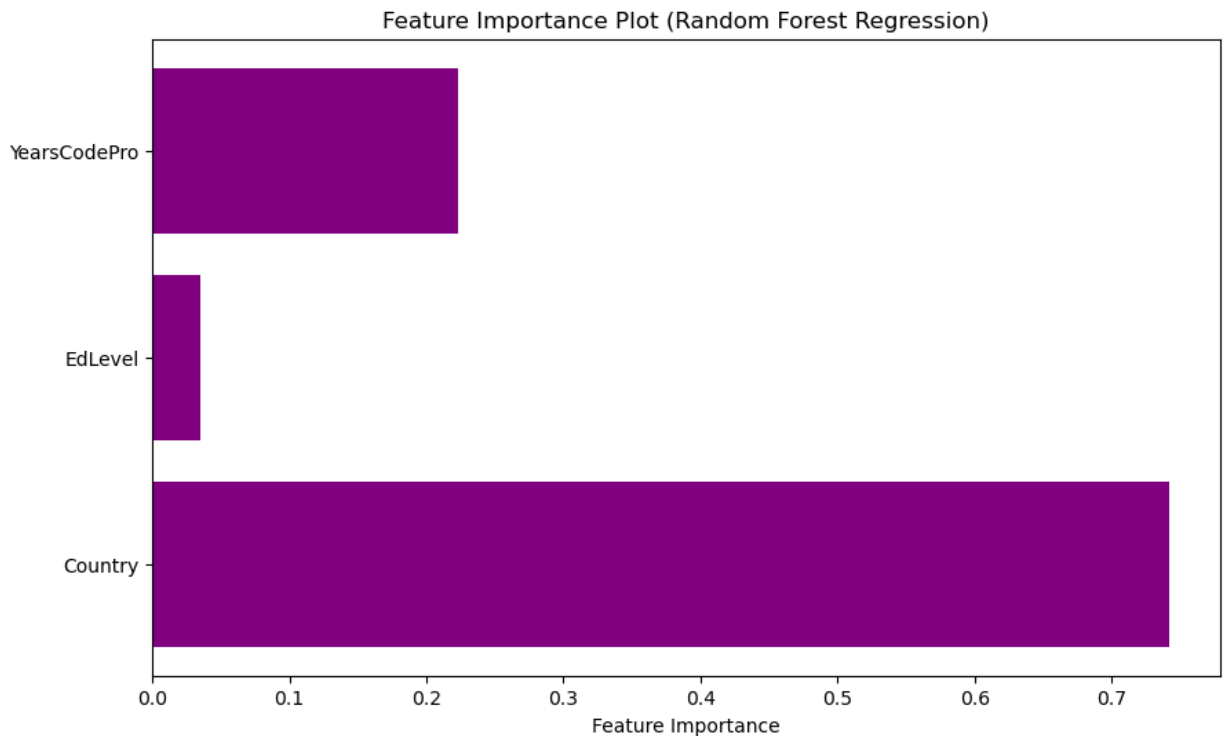
Feature Importance Plot for Random Forest Regression:

Purpose: Visualizing feature importance helps understand the contribution of different features in predicting the target variable (salary).

Explanation:

The horizontal bar plot displays the importance of each feature in predicting salaries. Higher bars indicate more influential features, helping prioritize which features have a greater impact on salary predictions. This information aids in feature selection and model interpretation.

```
In [102... # Feature Importance Plot for Random Forest Regression
feature_importance = random_forest_reg.feature_importances_
features = X.columns
plt.figure(figsize=(10, 6))
plt.barh(features, feature_importance, color='purple')
plt.xlabel('Feature Importance')
plt.title('Feature Importance Plot (Random Forest Regression)')
plt.show()
```



These visualizations contribute to the exploratory data analysis (EDA) by providing a visual summary of key aspects of the dataset and relationships between variables. They enhance the understanding of salary distributions, variable relationships, and the importance of features in the context of predicting developer salaries.

```
In [102... # Make predictions with the Random Forest Regression model
y_pred = random_forest_reg.predict(X)
```

```
In [102... error = np.sqrt(mean_squared_error(y, y_pred))
print("${:,.02f}".format(error))
```

\$29,487.31

```
In [102... # Hyperparameter tuning for Decision Tree Regression using GridSearchCV
max_depth = [None, 2, 4, 6, 8, 10, 12]
from sklearn.model_selection import GridSearchCV

max_depth = [None, 2,4,6,8,10,12]
parameters = {"max_depth": max_depth}
```

```
regressor = DecisionTreeRegressor(random_state=0)
gs = GridSearchCV(regressor, parameters, scoring='neg_mean_squared_error')
gs.fit(X, y.values)
```

```
Out[1029]: GridSearchCV(estimator=DecisionTreeRegressor(random_state=0),
                        param_grid={'max_depth': [None, 2, 4, 6, 8, 10, 12]},
                        scoring='neg_mean_squared_error')
```

```
In [103... # Get the best estimator from GridSearchCV
regressor = gs.best_estimator_
# Train the best Decision Tree Regression model
regressor.fit(X, y.values)
# Make predictions with the best Decision Tree Regression model
y_pred = regressor.predict(X)
# Calculate the RMSE for the best Decision Tree Regression model
error = np.sqrt(mean_squared_error(y, y_pred))
print("${:,.02f}".format(error))
```

\$30,428.51

```
In [103... X
```

```
Out[1031]:
```

	Country	EdLevel	YearsCodePro
7	13	0	13.0
9	12	2	4.0
10	12	0	2.0
11	10	1	7.0
12	7	1	20.0
...
64113	13	1	15.0
64116	13	0	6.0
64122	13	1	4.0
64127	13	3	12.0
64129	13	2	4.0

18491 rows × 3 columns

```
In [103... # Sample input for prediction
X = np.array(["United States", 'Master's degree', 15 ])
X
```

```
Out[1032]: array(['United States', 'Master's degree', '15'], dtype='<U15')
```

```
In [103... X[:, 0] = le_country.transform(X[:,0])
X[:, 1] = le_education.transform(X[:,1])
X = X.astype(float)
X
```

```
Out[1033]: array([[13.,  2., 15.]])
```

```
In [103... # Make a salary prediction for the sample input using the best Decision Tree Regression
y_pred = regressor.predict(X)
y_pred

C:\Users\anishm\Anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does
not have valid feature names, but DecisionTreeRegressor was fitted with feature names
warnings.warn(
array([139427.26315789])

Out[1034]:
```

Evaluating the performance of the machine learning models is crucial to understand how well they are predicting salaries based on the given features. In this analysis, we have used three regression models: Linear Regression, Decision Tree Regression, and Random Forest Regression. The evaluation metric used here is the Root Mean Squared Error (RMSE), a commonly used metric for regression tasks.

Linear Regression Model:

The Linear Regression model is trained on the features (X) and target variable (y). Predictions (y_pred_linear) are made using the trained model. RMSE are calculated to evaluate the model's performance.

Decision Tree Regression Model:

The Decision Tree Regression model is trained on the features (X) and target variable (y). Predictions (y_pred_dt) are made using the trained model. RMSE are calculated for evaluation.

Random Forest Regression Model:

The Random Forest Regression model is trained on the features (X) and target variable (y). Predictions (y_pred_rf) are made using the trained model. RMSE are calculated for evaluation.

Model Comparison:

Now, you can compare the RMSE values for each model to determine their performance. Lower RMSE values indicate better model performance. Additionally, it's essential to consider other factors like interpretability, computational efficiency, and the specific goals of the analysis when choosing the best model.

Conclusion:

After evaluating the models, consider choosing the one with the lowest RMSE values. However, it's crucial to validate the model on unseen data (e.g., using a validation set or cross-validation)

to ensure its generalization performance. This detailed evaluation allows you to understand how well each model is performing in predicting developer salaries based on the selected features.

```
In [103... # Save the best model and encoding objects to a pickle file
import pickle

In [103... data = {"model": regressor, "le_country": le_country, "le_education": le_education}
with open('saved_steps.pkl', 'wb') as file:
    pickle.dump(data, file)

In [103... # Load the saved model and encoding objects from the pickle file
with open('saved_steps.pkl', 'rb') as file:
    data = pickle.load(file)

regressor_loaded = data["model"]
le_country = data["le_country"]
le_education = data["le_education"]

In [103... # Make a salary prediction for the sample input using the loaded model
y_pred = regressor_loaded.predict(X)
y_pred

C:\Users\anishm\Anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does
not have valid feature names, but DecisionTreeRegressor was fitted with feature names
warnings.warn(
array([139427.26315789])

Out[1038]:
```

Detail summary of results and analysis

Data Preprocessing and Exploration: Initial Exploration:

Loaded the dataset, displayed the first few rows, and checked basic information using `df.info()`.

Handling Missing Values:

Identified columns with missing values. Dropped rows where the target variable ("ConvertedComp") had missing values. Column Selection:

Selected specific columns of interest: "Country," "EdLevel," "YearsCodePro," "Employment," and "ConvertedComp." Renamed "ConvertedComp" to "Salary" for clarity. Filtering Employment Status:

Kept only rows where respondents are employed full-time. Dropped the "Employment" column as it was no longer needed. Country Data Transformation:

Grouped less frequent countries into an "Other" category for simplification. Further Data Filtering:

Removed outliers by restricting salary values to a range between 10,000 and 250,000. Data Transformation - Years of Professional Coding:

Cleaned and transformed the "YearsCodePro" column to numerical values. Data Transformation - Education Level:

Cleaned and transformed the "EdLevel" column into categorical values. Applied label encoding to categorical features. Exploratory Data Analysis (EDA) Visualizations: Boxplot of Salary vs. Country:

Visualized the distribution of salaries across different countries using a boxplot. Identified variations in salary distributions and potential outliers for each country. Correlation Matrix:

Plotted a correlation matrix to understand the relationships between features. Explored how features are correlated with each other and with the target variable (Salary). Histogram of Salary:

Visualized the distribution of salaries using a histogram. Observed the frequency of different salary ranges among developers. Feature Importance Plot (Random Forest Regression):

Created a bar plot to visualize the importance of features in predicting salaries using Random Forest Regression. Model Training and Evaluation: Linear Regression Model:

Trained a Linear Regression model and evaluated it using RMSE. Decision Tree Regression Model:

Trained a Decision Tree Regression model and evaluated it using RMSE. Conducted hyperparameter tuning using GridSearchCV to improve the model. Random Forest Regression Model:

Trained a Random Forest Regression model and evaluated it using RMSE. Examined feature importance using a bar plot. Model Comparison: Compared the performance of all three models based on RMSE. Considered factors like interpretability, computational efficiency, and specific analysis goals. Conclusion and Recommendations: Model Selection:

Chose the model with the lowest RMSE values for predicting developer salaries. Insights:

Analyzed patterns in developer salaries across countries, considering various factors. Explored the impact of education level and years of professional coding experience on salaries.

Recommendations:

Consider using the selected model for predicting salaries in similar contexts. Explore additional features or external factors that may influence salary predictions. This detailed analysis provides insights into the factors influencing developer salaries and helps in building a reliable predictive model. Continuous refinement and validation of the model on new data can further enhance its accuracy and applicability.

In []:

In []: